



Machine Learning Engineer (NLP)

San Francisco Bay Area, New York City, or Remote

Revelata, Inc. is a Palo Alto-based, early-stage startup on a mission to unlock the actionable information held in *any* financial document. Our machine learning platform automates the drudgery out of reading reams of text, just for that one valuable nugget. With our data service, finance professionals can bridge the chasm from text to insight and make better-informed decisions – *faster*.

Natural language processing is central to our mission. We use it to efficiently discover, extract, and synthesize data from unstructured documents, replacing countless man-hours of toil with a lightweight data service, where every data point is up-to-date, accurate, and easily traced back to its source.

We are looking for a founding **Machine Learning Engineer** to join us in building a world-class ML platform that helps our clients make use of the information embedded in the billions of financial documents found worldwide and across all industries. As a founding engineer, you will take ownership and build significant portions of our NLP pipeline and work over the entire ML development lifecycle, from prototyping to deployment. You will also play a central role in shaping our ML team as we expand.

The exceptional candidate would have a track record of successfully training and productizing NLP pipelines to solve complex, real-world problems at scale, or a strong publication history in deep learning/NLP or application areas. We are particularly interested in candidates with a background in active learning and weak supervision, and who thrive in a scrappy, small-team environment.

Our founding team and advisors are all Harvard alumni, with decades of combined experience in machine learning and private equity/banking.

Qualifications:

- M.S. or Ph.D. in Computer Science or related field, with a focus on machine learning or NLP. Recent Harvard College (A.B.) alumni with industry/internship experience and relevant upper-level coursework in Computer Science are also welcome to apply.
- 3-5 years of experience in ML/NLP:
 - Expert in deep learning and/or neural language models
 - Kept up with the latest industry and academic research
 - Experienced in (and like) wrangling messy data and scoping complex problems
 - Have customized/trained/fine-tuned Transformer or other neural language models
 - Shipped code in ML products
 - Practice good MLOps habits
 - Experience with Python and AWS
- Care deeply about software engineering fundamentals, including clear documentation
- Care deeply about strong business outcomes
- Ability to work effectively by yourself and in a team
- Open communication style

Contact: CK Lin <ck@revelata.com>